

Parallel Matrix Multiplication

Final Project

INF236: Parallel Programming

**Kateřina Čížková, Luca
Klingenberg**

UNIVERSITY OF BERGEN



Table of contents

1 Introduction

2 Strassen's Algorithm

- Different values for the cutoff level
- Z-ordering

3 Parallelization

- Parallel matrix multiplication
- Parallel Strassen

4 Experiments

- Comparing Strassen algorithm variants
- Matrices of different sizes

5 Conclusion



Introduction

- Complexity: $\mathcal{O}(n^3)$
- Consecutive memory access

Algorithm 1 matrix multiplication

Input: **A**, **B**

Output: **C** (the resulting matrix)

```
1: function MATMUL(A, B)
2:   for  $i = 0, \dots, n - 1$  do
3:     for  $j = 0, \dots, n - 1$  do
4:        $c[i][j] = 0$ 
5:       for  $k = 0, \dots, n - 1$  do
6:         for  $j = 0, \dots, n - 1$  do
7:            $c[i][j] += a[i][k] \cdot b[k][j]$ 
8:   return C
```



Table of contents

1 Introduction

2 Strassen's Algorithm

- Different values for the cutoff level
- Z-ordering

3 Parallelization

- Parallel matrix multiplication
- Parallel Strassen

4 Experiments

- Comparing Strassen algorithm variants
- Matrices of different sizes

5 Conclusion



Strassen's Algorithm

- Blockwise matrix multiplication

$$\begin{aligned}\mathbf{A} \cdot \mathbf{B} &= \begin{pmatrix} \mathbf{A}_{00} & \mathbf{A}_{01} \\ \mathbf{A}_{10} & \mathbf{A}_{11} \end{pmatrix} \cdot \begin{pmatrix} \mathbf{B}_{00} & \mathbf{B}_{01} \\ \mathbf{B}_{10} & \mathbf{B}_{11} \end{pmatrix} \\ &= \begin{pmatrix} \mathbf{A}_{00}\mathbf{B}_{00} + \mathbf{A}_{01}\mathbf{B}_{10} & \mathbf{A}_{00}\mathbf{B}_{01} + \mathbf{A}_{01}\mathbf{B}_{11} \\ \mathbf{A}_{10}\mathbf{B}_{00} + \mathbf{A}_{11}\mathbf{B}_{10} & \mathbf{A}_{10}\mathbf{B}_{01} + \mathbf{A}_{11}\mathbf{B}_{11} \end{pmatrix} \\ &= \begin{pmatrix} \mathbf{C}_{00} & \mathbf{C}_{01} \\ \mathbf{C}_{10} & \mathbf{C}_{11} \end{pmatrix} = \mathbf{C}\end{aligned}$$

- Idea: reduce number of multiplications from 8 to 7
- Complexity of Strassen's algorithm: $\mathcal{O}(n^{\log_2 7}) \approx \mathcal{O}(n^{2.8073})$
- Can be further improved by reusing intermediate results of additions and subtractions (Winograd's algorithm)
- Input matrices are zero-padded until the next power of two



Strassen's Algorithm

Algorithm 2 Strassen's matrix multiplication

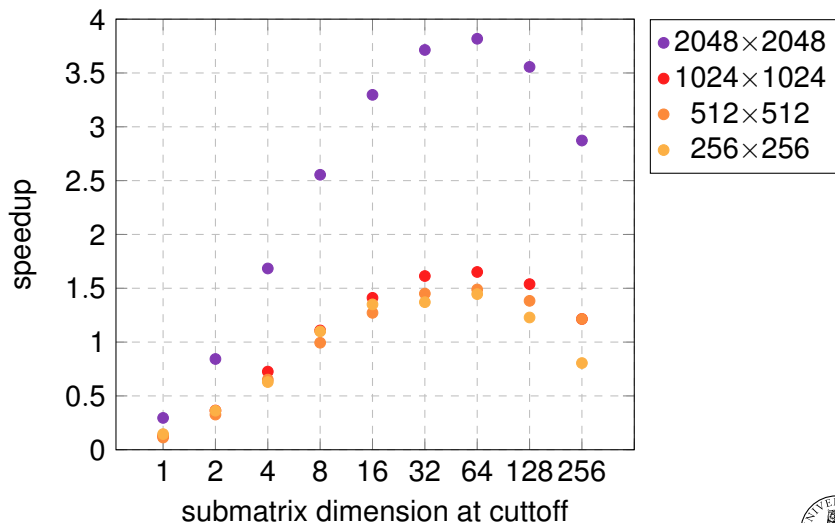
Input: A, B

Output: C (the resulting matrix)


```
1: function STRASSEN( $A, B, n$ )
2:   if  $n == \text{cutoff}$  then
3:     return MATMUL( $A, B$ )
4:    $P_1 = \text{STRASSEN}(A_{00} + A_{11}, B_{00} + B_{11}, \frac{n}{2})$ 
5:    $P_2 = \text{STRASSEN}(A_{10} + A_{11}, B_{00}, \frac{n}{2})$ 
6:    $P_3 = \text{STRASSEN}(A_{00}, B_{01} - B_{11}, \frac{n}{2})$ 
7:    $P_4 = \text{STRASSEN}(A_{11}, B_{10} - B_{00}, \frac{n}{2})$ 
8:    $P_5 = \text{STRASSEN}(A_{00} + A_{01}, B_{11}, \frac{n}{2})$ 
9:    $P_6 = \text{STRASSEN}(A_{10} - A_{00}, B_{00} + B_{01}, \frac{n}{2})$ 
10:   $P_7 = \text{STRASSEN}(A_{01} - A_{11}, B_{10} + B_{11}, \frac{n}{2})$ 
11:   $C_{00} = P_1 + P_4 - P_5 + P_7$ 
12:   $C_{01} = P_3 + P_5$ 
13:   $C_{10} = P_2 + P_4$ 
14:   $C_{11} = P_1 - P_2 + P_3 + P_6$ 
15:  return  $C$ 
```



Different values for the cutoff level




Z-ordering




a_{00} a_{01} a_{02} a_{03}
 a_{10} a_{11} a_{12} a_{13}
 a_{20} a_{21} a_{22} a_{23}
 a_{30} a_{31} a_{32} a_{33}

a_{00} a_{01} a_{02} a_{03}
 a_{10} a_{11} a_{12} a_{13}
 a_{20} a_{21} a_{22} a_{23}
 a_{30} a_{31} a_{32} a_{33}



1D array in memory:

a_{00} a_{01} a_{02} a_{03} a_{10} a_{11} a_{12} ...




1D array in memory:

a_{00} a_{01} a_{10} a_{11} a_{02} a_{03} a_{12} ...



Z-ordering

- Matrices at cutoff level are stored row-wise



- Only pointers instead of the actual submatrices are passed to the recursive calls



Table of contents

1 Introduction

2 Strassen's Algorithm

- Different values for the cutoff level
- Z-ordering

3 Parallelization

- Parallel matrix multiplication
- Parallel Strassen

4 Experiments

- Comparing Strassen algorithm variants
- Matrices of different sizes

5 Conclusion



Parallelization: Matrix multiplication

Algorithm 3 matrix multiplication

Input: A, B

Output: C (the resulting matrix)

```
1: #pragma omp parallel for
2: function MATMUL( $A, B$ )
3:   for  $i = 0, \dots, n - 1$  do
4:     for  $j = 0, \dots, n - 1$  do
5:        $c[i][j] = 0$ 
6:       for  $k = 0, \dots, n - 1$  do
7:         for  $j = 0, \dots, n - 1$  do
8:            $c[i][j] += a[i][k] \cdot b[k][j]$ 
9:   return  $C$ 
```



Parallelization: Strassen

■ 2-layers variant

- Use 2 hardcoded recursion levels of Strassen
- Execute all additions and subtractions in parallel
- At the cutoff level, execute all matrix multiplications in parallel

■ Recursive variant

- Use several recursive calls of Strassen
- Execute all additions and subtractions in parallel
- At the cutoff level, execute all matrix multiplications in parallel
- Strassen is recursively applied until the submatrices have a dimension of 512×512



Table of contents

1 Introduction

2 Strassen's Algorithm

- Different values for the cutoff level
- Z-ordering

3 Parallelization

- Parallel matrix multiplication
- Parallel Strassen

4 Experiments


- Comparing Strassen algorithm variants
- Matrices of different sizes

5 Conclusion



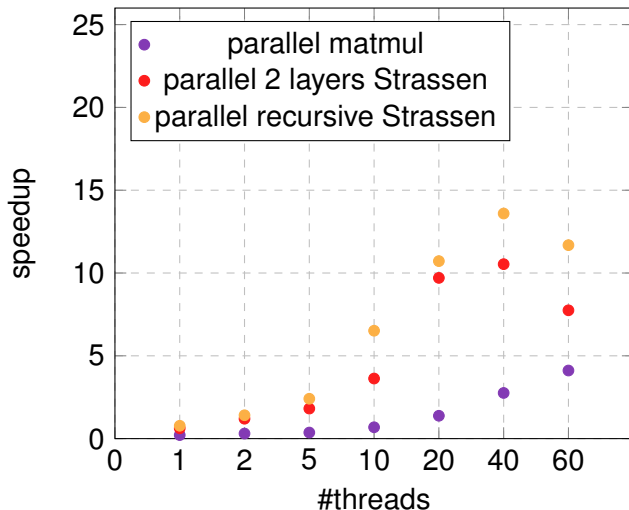
Experiments

Comparing Strassen algorithm variants (256×256)




Experiments

Comparing Strassen algorithm variants (4096×4096)



Experiments

Comparing Strassen algorithm variants (8192×8192)



Experiments

Running time for matrices of different sizes

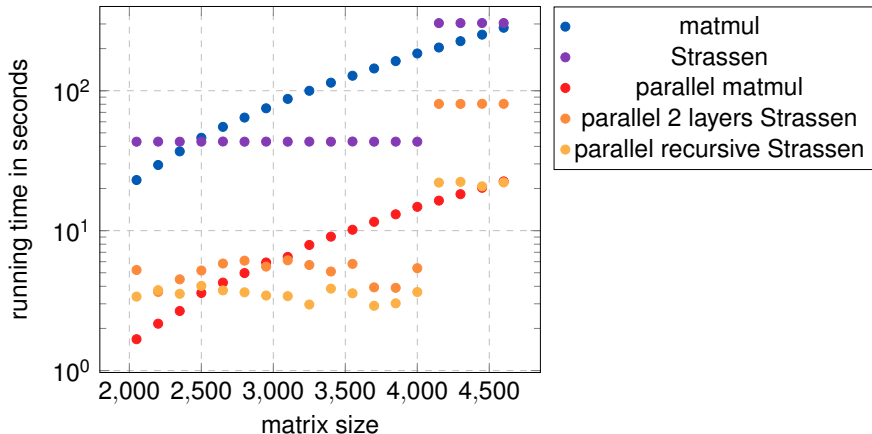


Table of contents

1 Introduction

2 Strassen's Algorithm

- Different values for the cutoff level
- Z-ordering

3 Parallelization

- Parallel matrix multiplication
- Parallel Strassen

4 Experiments

- Comparing Strassen algorithm variants
- Matrices of different sizes

5 Conclusion



Conclusion

- With sufficient matrix size, the parallel variants of Strassen's algorithm perform much better
- For more than 40 threads, the parallel recursive Strassen version slows down again
- Zero padding could be improved

